

CS 553 Cloud Computing Homework 5

➤ Problem

Implement the Shared-Memory TeraSort application for four datasets of different sizes : 1GB, 4GB, 16GB, 64GB. Memory usage is limited to 8 GB.

➤ Methodology

- For 1GB and 4 GB of dataset we use in- memory sort
- For 16 GB and 64 GB we use external sort
- We use parallel merge sort for all four data sets.
- And use k -way merge for external sort we combine the result into a single file.

➤ Runtime Environment

- Compute node: Skylake node
- C++
- Linux
- Memory : 8 GB

➤ Overall benchmark design

- MySort.cpp is the main file
- Input to the file includes no. of threads and dataset file.
 - ./MySort 24 record_1gb.txt
- getFileSize() function is used to get the size of the file on bytes.
- If the size of the file is less than 8GB we call internal sort else, we perform external sort.
- For in-memory sort
 - First we read the data from the file and store it in a vector using readData() function
 - Once the data is read we pass the vector and the thread count to inplace_merge_sort() which performs parallel merge sort.
 - inplace_merge_sort() : we divide the dataset among threads depending on the no. of threads given by the user.
 - Each thread performs merge sort on its own data and returns the result.
 - Then we merge the result and the final sorted data is written into the output file using writeData().
- For external sort (Max Memory Size = 8GB)
 - We decide the no. of chunks the file has to be divided into using the formula:
 - $\text{Chunks} = \text{FileSize} / \text{Max Memory Size}$
 - Once the chunks are decided then we calculate the read and write buffer size.
 - $\text{Read buffer} = \text{Max Memory Size} / \text{Chunks} + 1$
 - $\text{Write buffer} = \text{Max Memory Size} - (\text{Read Buffer} * \text{Chunks})$
 - “counter” variable and memory_utilization_check() function keeps track of the memory being allocated and deallocated to make sure that we don’t exceed the given size.
 - Threads are created equal to the no. of chunks.
 - Each thread will read the max size allowed from the file and write it to another file.

- Once the file is divided into chunks, we sequentially sort each chunk using inplace merge sort.
- After all the files are sorted, we call the kway merge to create the final sorted file.
- For k-way merge, MinHeap is used to store min element of each file.

```

struct MinHeapNode
{
    string line;
    int i;
};

```
- All the input files are opened, then we store the min element from each file into array of MinHeap.
- Priority queue is used to compare the elements in min heap and store them in ascending order
- Then we pop the top element from the queue and write it to the file.
- Then we continue reading the next min element from all the files and store it queue.
- We perform the above steps until all the data from the files are read.
- The final output file is created with sorted data.


➤ Performance evaluation of Single Node TeraSort (using best # of threads for each case)

Experiment	Shared Memory (1GB)	Linux Sort (1GB)	Shared Memory (4GB)	Linux Sort (4GB)	Shared Memory (16GB)	Linux Sort (16GB)	Shared Memory (64GB)	Linux Sort (64GB)
Number of Threads	24	24	24	24	48+2(IO)	48	48+8(IO)	48
Sort Approach (e.g. in-memory / external)	In-memory	External	In-memory	External	External	External	External	External
Sort Algorithm (e.g. quicksort / mergesort / etc)	Merge-sort	Merge-sort	Merge-sort	Merge-sort	Merge-sort	Merge-sort	Merge-sort	Merge-sort
Data Read (GB)	1	1	4	4	3	2.58	0.89	2.58
Data Write (GB)	1	1	4	4	2	2.58	0.89	2.58
Sort Time (sec)	87.21	50.03	366.796	315.89	2138.51	995.92	8469.04	5328.08
Overall I/O Throughput (MB/sec)	11.74	20.46	11.167	12.97	7.66	16.45	7.74	12.3
Overall CPU Utilization (%)	88.24	55.6	92.3	61.4	99.6	98.7	100	99.7
Average Memory Utilization (GB)	0.13	0.13	1.71	1.14	4.08	1.20	4.27	1.21


➤ Results

- For 1GB and 4GB dataset MySort performs in-memory merge sort with throughput of 11.74 and 11.16 respectively. Whereas Linux sort has a throughput of 20.46 and 12.97 respectively.
- Sort time and Average memory utilization is similar for both MySort and Linux Sort for 1GB and 4 GB dataset.
- For 16 GB and 64 GB dataset MySort takes more time as compared to Linux sort and throughput is also less. MySort uses external merge sort along with k-way merge to sort the data.
- Average memory utilization is greater for MySort as compared to Linus Sort for 16GB and 64 GB dataset.


➤ Graph for Dataset vs Sort time for both sorts



➤ MySort 64GB graph



➤ Linux Sort 64GB graph



- For MySort, CPU utilization spikes when sorting each chunk of the file and merging the files.
- The memory utilization is also high compared to Linux Sort.

➤ Logs

- **MySort 1GB**

Data Read (GB): 1

Data Write (GB): 1

Sort Time (sec): 87.21

Throughput(MB/sec): 11.74

./valsort record_1GB_sorted.txt

Records: 10737418

Checksum: 51e82f1348dc5f

Duplicate keys: 335752

SUCCESS - all records are in order

- **Linux Sort 1GB**

Data Read (GB): 1

Data Write (GB): 1

Sort Time (sec): 50.03

Throughput(MB/sec): 20.46

./valsort output_1GB.txt

Records: 10737418
Checksum: 51e88faefe7b34
Duplicate keys: 0
SUCCESS - all records are in order

- **MySort 4GB**

Data Read (GB): 4

Data Write (GB): 4

Sort Time (sec): 366.796

Throughput(MB/sec): 11.167

./valsort record_4GB_sorted.txt
Records: 42949672
Checksum: 147a6e74151eb00
Duplicate keys: 2426680
SUCCESS - all records are in order

- **Linux Sort 4GB**

Data Read (GB): 4

Data Write (GB): 4

Sort Time (sec): 315.89

Throughput(MB/sec): 12.97

./valsort output_4GB.txt
Records: 42949672
Checksum: 147a91dd0ea4d36
Duplicate keys: 0
SUCCESS - all records are in order

- **MySort 16GB**

Data Read (GB): 3

Data Write (GB): 2

Sort Time (sec): 2138.51

Throughput(MB/sec): 7.66

./valsort output_sorted.txt
Records: 171798690

Checksum: 51ecbbabb0942db
Duplicate keys: 86040842
SUCCESS - all records are in order

- **Linux Sort 16GB**

Data Read (GB): 2.58

Data Write (GB): 2.58

Sort Time (sec): 995.92

Throughput(MB/sec): 16.45

./valsort output_16GB.txt
Records: 171798691
Checksum: 51ea5e584e0bcd3
Duplicate keys: 0
SUCCESS - all records are in order

- **MySort 64GB**

Data Read (GB): 0.89

Data Write (GB): 0.89

Sort Time (sec): 8469.04

Throughput(MB/sec): 7.74

./valsort output_sorted.txt
Records: 687194760
Checksum: 147abb9423e3f49d
Duplicate keys: 601295497
SUCCESS - all records are in order

- **Linux Sort 64GB**

Data Read (GB): 2.58

Data Write (GB): 2.58

Sort Time (sec): 5328.08

Throughput(MB/sec): 12.3

./valsort output_64GB.txt
Records: 687194767

Checksum: 147ad98a21f6e7a7
Duplicate keys: 0
SUCCESS - all records are in order